



## ucvm2etree


Map projection from a longitude-latitude-depth to a x-y-z coordinate system.




In the single-core command, the model domain is divided into  $c_x \times c_y$  columns. Each column is meshed as an independent octree.





As ucvm2etree queries the meta-model, it stores the data-point payloads into the etree using `etree_insert()` from the etree library. Since the inserts are not done in global in z-order, the outcome does not optimize disk-space.



A recommended step after running ucvm2etree is to run program `ecomact`. This code traverses the etree generated by ucvm2etree and builds a copy by appending octants in z-order. The outcome is an equivalent smaller file, optimal for querying performance.




The end result is a binary file (etree) with metadata about the model origin coordinates, dimensions, date of creation and authorship.




## ucvm2etree\_[mpi-process]


In the first step of the parallel commands, ucvm2etree\_extract maps the domain to  $p_x \times p_y$  processors.




Each processor receives  $c_x/p_x \times c_y/p_y$  columns and meshes each column independently.



In ucvm2etree\_extract, each processor queries the meta-model by columns and stores the data-points in column meshes. The octants in the mesh of each column are arranged in **local z-order** and written on disk.



In ucvm2etree\_sort  
The local column meshes are sorted in **global z-order** so they can later be merged, but remain in separate files on disk.



The last step the parallel version, ucvm2etree\_merge, merges the global z-ordered column files into a single mesh.

